


Importing Libraries

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import cv2
from google.colab.patches import cv2_imshow
from PIL import Image
import tensorflow as tf
tf.random.set_seed(3)
from tensorflow import keras
from keras.datasets import mnist
from tensorflow.math import confusion_matrix
```

Loading the MNIST data from keras.datasets


```
(X_train, Y_train), (X_test, Y_test) = mnist.load_data()
```

 Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>


11490434/11490434

0s 0us/step

```
type(X_train)
```

 numpy.ndarray

```
# Shape of the numpy arrays
print(X_train.shape, Y_train.shape, X_test.shape, Y_test.shape)
```

 (60000, 28, 28) (60000,) (10000, 28, 28) (10000,)

Above output shows that


Training data = 60,000 Images

Test data = 10,000 Images

Image dimension --> 28x28

Grayscale Image --> 1 Channel

```
# Pinting the 10th Image
print(X_train[10])
```



```
[[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0 42 118 219 166 118 118    6
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0 103 242 254 254 254 254   66
   0  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0 18 232 254 254 254 254 238
  70  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0 104 244 254 224 254 254
 141  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 207 254 210 254 254 254
 34  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 84 206 254 254 254 254
 41  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 24 209 254 254 254
 171  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 91 137 253 254 254 254
 112  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0 40 214 250 254 254 254 254 254
 34  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0 81 247 254 254 254 254 254 254
 146  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0 110 246 254 254 254 254 254
 171  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0 73 89 89 93 240 254
 171  0  0  0  0  0  0  0  0  0]
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  1 128 254
219 31  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  7 254 254
214 28  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 138 254 254
116  0  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0 19 177 90  0  0  0  0  0 25 240 254 254
34  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0 164 254 215 63 36  0 51 89 206 254 254 139
8  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0 57 197 254 254 222 180 241 254 254 253 213 11
0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0 140 105 254 254 254 254 254 236  0  0
0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  7 117 117 165 254 254 239 50  0  0
0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0  0  0  0  0]]
```

```
print(X_train[10].shape)
```

➡ (28, 28)

```
# Displaying the image
plt.imshow(X_train[25])
plt.show()
# printing the corresponding image
print(Y_train[50])
```

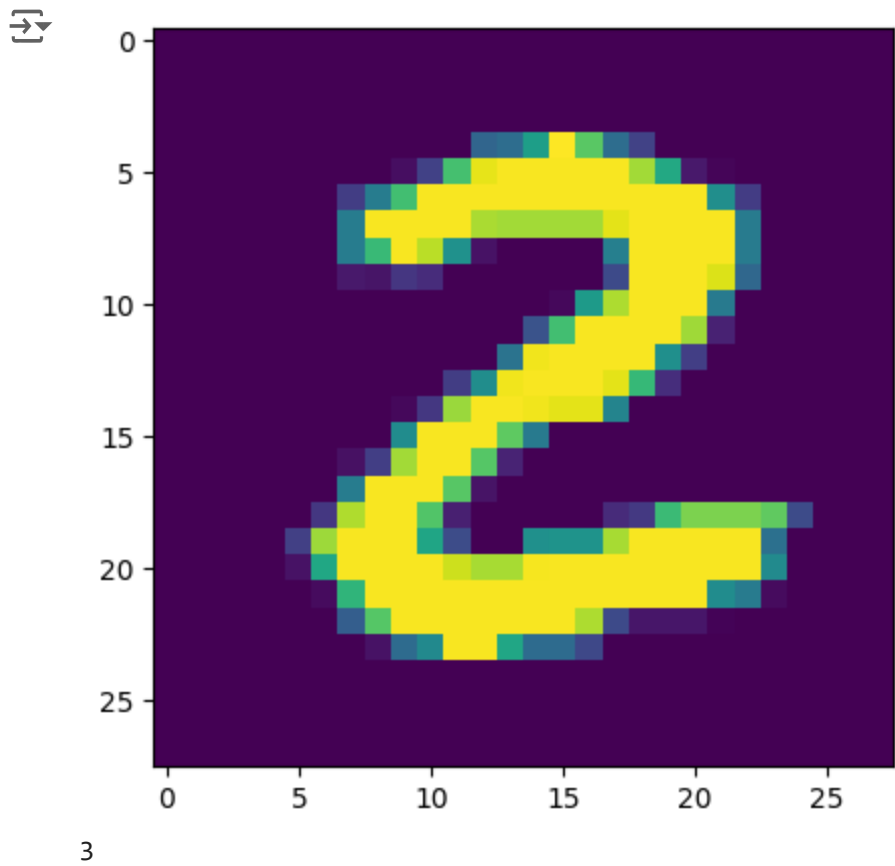


Image Lables

```
print(Y_train.shape, Y_test.shape)
```

➡ (60000,) (10000,)

```
# Unique Value in Y_train
print(np.unique(Y_train))

# Unique Value in Y_test
print(np.unique(Y_test))
```

➡ [0 1 2 3 4 5 6 7 8 9]
[0 1 2 3 4 5 6 7 8 9]

```
# Scaling the values
X_train = X_train/255
X_test = X_test/255
```

```
# Printing the 20th image
print(X_train[20])
```

```

0.19215686 0.98823529 0.98823529 0.98823529 0.98823529 0.98823529
0.98823529 0.85098039 0.84705882 0.55294118 0.49411765 0.98823529
0.98823529 0.98823529 0.60784314 0. 0. 0.
0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0.
0.19215686 0.98823529 0.98823529 0.98823529 0.91764706 0.8
0.34901961 0. 0. 0. 0. 0.19215686 0.98823529
0.98823529 0.98823529 0.60784314 0. 0. 0.
0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0.
0.05490196 0.61960784 0.75294118 0.59215686 0.17647059 0.
0. 0. 0. 0. 0.19215686 0.98823529
0.98823529 0.98823529 0.88235294 0.06666667 0. 0.
0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0.19215686 0.98823529
0.98823529 0.98823529 0.98823529 0.09019608 0. 0.
0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0.12941176 0.89411765
0.98823529 0.98823529 0.98823529 0.61568627 0.01568627 0.
0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.21568627
0.89803922 0.98823529 0.98823529 0.98823529 0.04313725 0.
0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0.20784314 0.90980392 0.98823529 0.98823529 0.24705882 0.
0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0.35294118 0.80784314 0.51372549 0.04313725 0.
0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. ]
[0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. ]

```

Building the Neural Network

```
# Setting up the layers of the Neural Network
model = keras.Sequential([
    keras.layers.Flatten(input_shape=(28,28)),
    keras.layers.Dense(50, activation='relu'),
    keras.layers.Dense(10, activation='sigmoid')
])
```

```
# Compiling the Neural Network
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
# Training the Neural Network
model.fit(X_train, Y_train, epochs=10)
```

↗ Epoch 1/10
1875/1875 ————— 5s 2ms/step - accuracy: 0.8478 - loss: 0.5343
Epoch 2/10
1875/1875 ————— 3s 2ms/step - accuracy: 0.9539 - loss: 0.1591
Epoch 3/10
1875/1875 ————— 3s 2ms/step - accuracy: 0.9671 - loss: 0.1123
Epoch 4/10
1875/1875 ————— 3s 2ms/step - accuracy: 0.9744 - loss: 0.0884
Epoch 5/10
1875/1875 ————— 4s 2ms/step - accuracy: 0.9791 - loss: 0.0729
Epoch 6/10
1875/1875 ————— 3s 2ms/step - accuracy: 0.9824 - loss: 0.0620
Epoch 7/10
1875/1875 ————— 3s 2ms/step - accuracy: 0.9848 - loss: 0.0532
Epoch 8/10
1875/1875 ————— 5s 2ms/step - accuracy: 0.9869 - loss: 0.0461
Epoch 9/10
1875/1875 ————— 4s 2ms/step - accuracy: 0.9887 - loss: 0.0401
Epoch 10/10
1875/1875 ————— 3s 2ms/step - accuracy: 0.9902 - loss: 0.0349
<keras.src.callbacks.history.History at 0x7cad02c74ac0>

Above output shows that Training data accuracy = 99.02%

Accuracy on Test Data:

```
loss, accuracy = model.evaluate(X_test, Y_test)
print(f"Test Loss: {loss}")
print(f"Test Accuracy: {accuracy}")
```

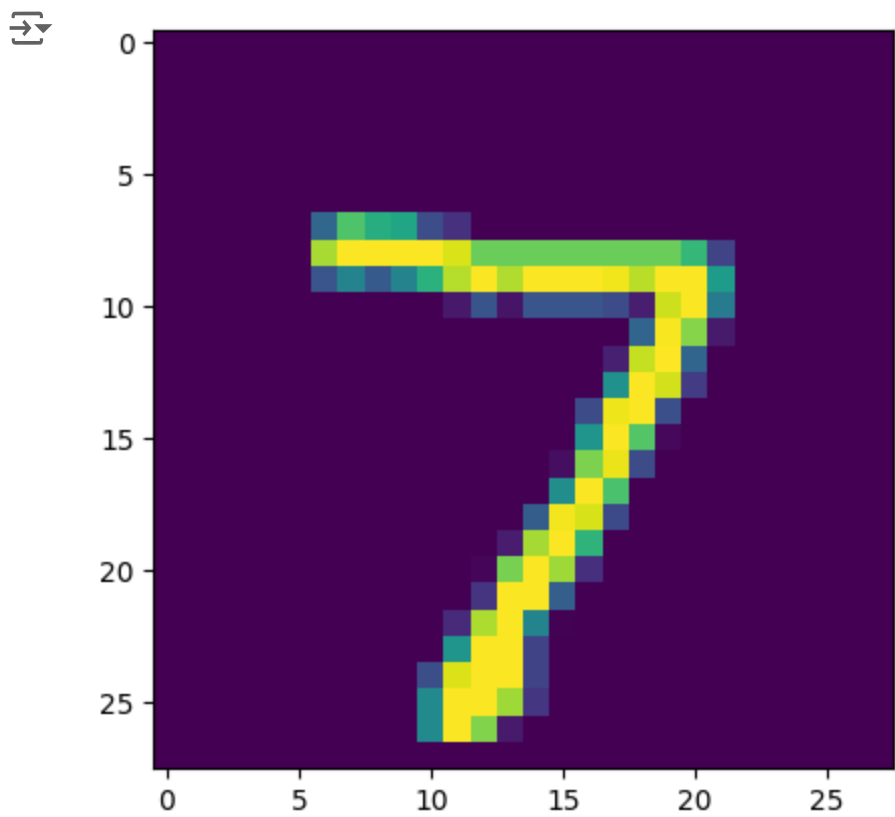
↗ 313/313 ————— 0s 1ms/step - accuracy: 0.9692 - loss: 0.1083
Test Loss: 0.09586985409259796
Test Accuracy: 0.9718999862670898

Test data accuracy = 97.18%

```
print(X_test.shape)
```

↗ (10000, 28, 28)

```
plt.imshow(X_test[0])
plt.show()
```



```
print(Y_test[0])
```

↗ 7

```
Y_pred = model.predict(X_test)
```

↗ 313/313 ————— 0s 1ms/step

```
print(Y_pred.shape)
```

(10000, 10)

```
print(Y_pred[0])
```

[6.6705686e-03 3.1692503e-05 6.7617349e-02 9.5477194e-01 9.2853525e-10
4.3258894e-02 3.2479630e-08 9.999344e-01 2.7314574e-01 6.6098881e-01]

```
label_for_first_test_image = np.argmax(Y_pred[0])  
print(label_for_first_test_image)
```

7

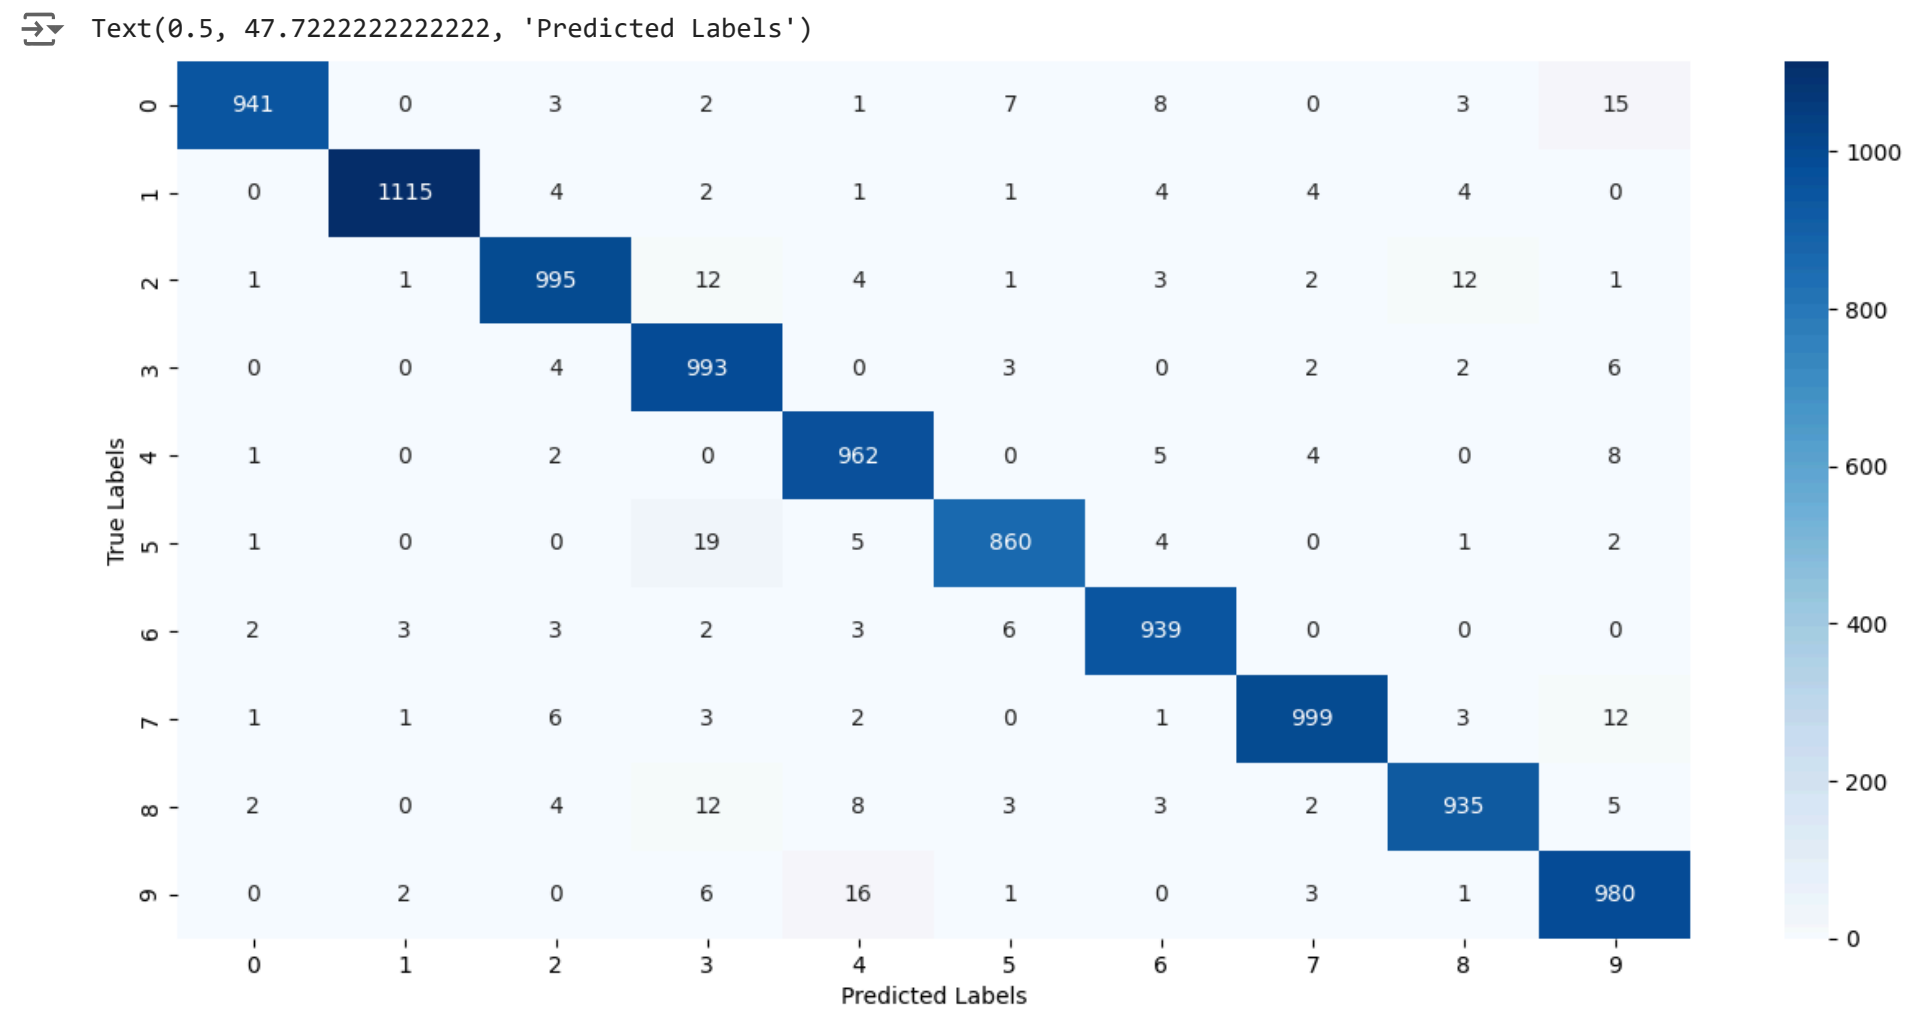
```
Y_pred_labels = [np.argmax(i) for i in Y_pred]  
print(Y_pred_labels)
```

[7, 2, 1, 0, 4, 1, 4, 9, 5, 9, 0, 6, 9, 0, 1, 5, 9, 7, 3, 4, 9, 6, 6, 5, 4, 0, 7, 4, 0, 1, 3, 1, 3, 4, 7, 2, 7, 1, 2, 1

```
conf_mat = confusion_matrix(Y_test, Y_pred_labels)  
print(conf_mat)
```

tf.Tensor(
[[941 0 3 2 1 7 8 0 3 15]
 [0 1115 4 2 1 1 4 4 4 0]
 [1 1 995 12 4 1 3 2 12 1]
 [0 0 4 993 0 3 0 2 2 6]
 [1 0 2 0 962 0 5 4 0 8]
 [1 0 0 19 5 860 4 0 1 2]
 [2 3 3 2 3 6 939 0 0 0]
 [1 1 6 3 2 0 1 999 3 12]
 [2 0 4 12 8 3 3 2 935 5]
 [0 2 0 6 16 1 0 3 1 980]], shape=(10, 10), dtype=int32)

```
plt.figure(figsize = (15,7))  
sns.heatmap(conf_mat, annot = True, fmt = 'd', cmap = 'Blues')  
plt.ylabel('True Labels')  
plt.xlabel('Predicted Labels')
```

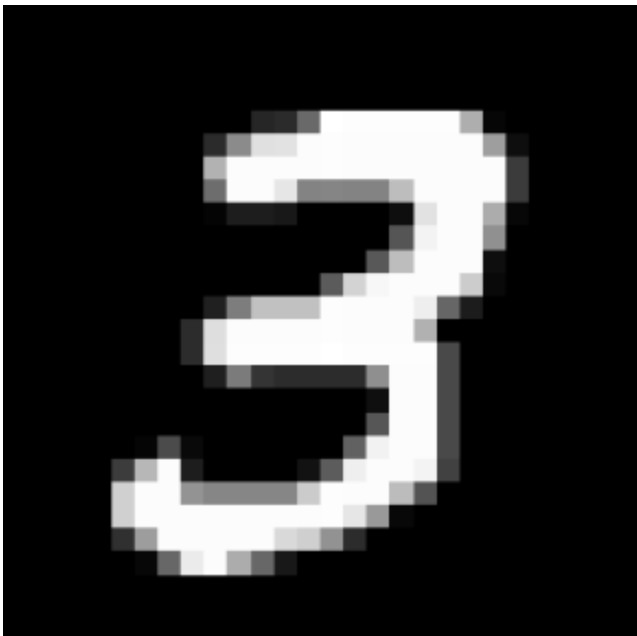


```
input_image_path = '/content/MNIST_digit.png'
input_image = cv2.imread(input_image_path)
```

```
type(input_image)
```

```
↔ numpy.ndarray
```

```
cv2.imshow(input_image)
```

```
↔ 
```

```
input_image.shape
```

```
↔ (318, 318, 3)
```

```
grayscale = cv2.cvtColor(input_image, cv2.COLOR_RGB2GRAY)
```

```
grayscale.shape
```

```
↔ (318, 318)
```

```
input_image_resize = cv2.resize(grayscale, (28,28))
```

```
input_image_resize.shape
```

```
↔ (28, 28)
```

```
cv2.imshow(input_image_resize)
```

```
↔ 
```

```
input_image_resize = input_image_resize/255
```

```
type(input_image_resize)
```

```
↔ numpy.ndarray
```

```
image_resaped = np.reshape(input_image_resize, [1,28,28])
```

```
input_prediction = model.predict(image_resaped)
print(input_prediction)
```

```
↔ 1/1 ————— 0s 21ms/step
[[1.4789014e-09 3.3444320e-10 2.8190700e-06 1.0000000e+00 1.5516093e-12
 2.0472713e-01 1.1959650e-18 1.8005638e-05 9.7836589e-04 6.8385339e-01]]
```

```
input_pred_label = np.argmax(input_prediction)
```

```
print(input_pred_label)
```


```
↔ 3
```

Predctive System

```
input_image_path = input('Path of the image to be predicted: ')
input_image = cv2.imread(input_image_path)
cv2_imshow(input_image)
grayscale = cv2.cvtColor(input_image, cv2.COLOR_RGB2GRAY)
input_image_resize = cv2.resize(grayscale, (28,28))
input_image_resize = input_image_resize/255
image_resaped = np.reshape(input_image_resize, [1,28,28])
input_prediction = model.predict(image_resaped)
input_pred_label = np.argmax(input_prediction)
print('The Handwritten Digit is Recognised as: ', input_pred_label)
```

 Path of the image to be predicted: /content/MNIST_digit.png



1/1  0s 19ms/step
The Handwritten Digit is Recognised as: 3